

## Informed Search Method

Use information about the state space to prevent algorithms from blundering about in the dark.

### Best-First Search

Uses the queuing function to order the nodes in the queue according to an **evaluation function**, and choose to expand a node that seems to have the best **desirability**.

**function** BEST-FIRST-SEARCH(*problem*, EVAL-FN) **returns** a solution sequence

**inputs:** *problem*, a problem

*Eval-Fn*, an evaluation function

*Queueing-Fn*  $\leftarrow$  a function that orders nodes by EVAL-FN

**return** GENERAL-SEARCH(*problem*, *Queueing-Fn*)

A common aspect of Best-First Search strategies is that, *in order to focus the search*, they use some ***estimated measure*** which incorporates some estimate of the cost of the path from a state to the closest goal

## Minimize estimated cost to reach a goal : Greedy search

The node whose state is judged to be the closest to the goal state is always expanded first.

✧ The strategy is to minimize the **remaining path cost**, and is called the **greedy search**.

A **heuristic function** is used to calculate the path cost:


$h(n)$  = estimated cost of the cheapest path  
from the state at node  $n$  to *a goal state*.

$h(n)$  must be **0** for goal states.

✧ The code is :

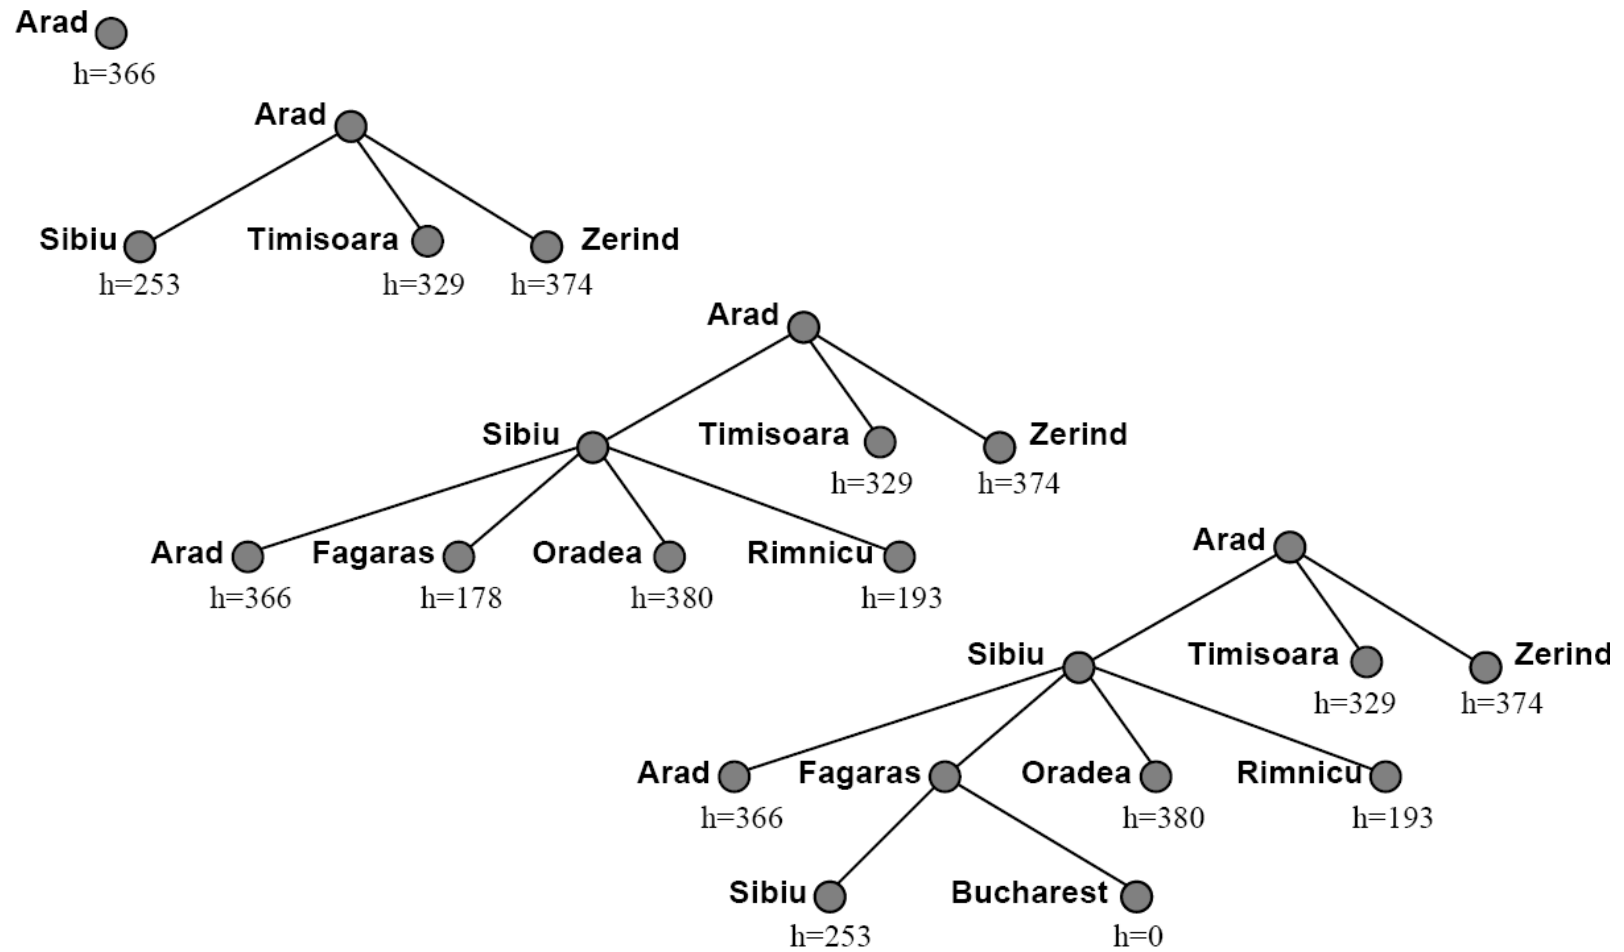
```
function GREEDY-SEARCH(problem)  
    returns a solution or failure.  
    return BEST-FIRST-SEARCH(problem,  $h$ )
```

✧ An example :



Define the heuristic function as:

$h_{SLD}(n)$  = straight-line distance between  $n$  and the *goal* location.



✧ Incomplete and not optimal:


1. greedy approaches may be **quick**. It may have the minimum search cost because of using *immediate best choice* instead of *long term options*.
2. greedy approaches may **not be optimal**. The solution found may not be the best one. ( Arad-Rimnicu-Vilcea-Pitesti-Bucharest is the best, for example.)
3. greedy approaches may **need to backtrack** (consider Iasi to Fagaras).

✧ Problem: holds all nodes in memory

✧ **time complexity == space complexity ==  $O(b^m)$ .**

## Minimize the total path cost : A\* search

Combines the Greedy search (quick) and Uniform-cost search (complete and optimal) to gain the benefits.



✧ The strategy is to minimize the **total path cost**.

An **evaluation function**  $f(n)$  is used to estimate the total path cost :

$$f(n) = \text{estimated cost of the cheapest solution through } n \\ = g(n) + h(n) \quad \text{where}$$

$g(n)$  = (known) cost of getting to  $n$

$h(n)$  = estimated cheapest cost to get from  $n$  to *goal*

In A\* search,  $h$  must be restricted to **admissible heuristic**, which is *inherently optimistic*, and *never overestimates the cost to reach the goal*.

The optimism is transferred to the  $f$  function : *if  $h$  is admissible,  $f(n)$  never overestimates the actual cost of the best solution through  $n$ .*

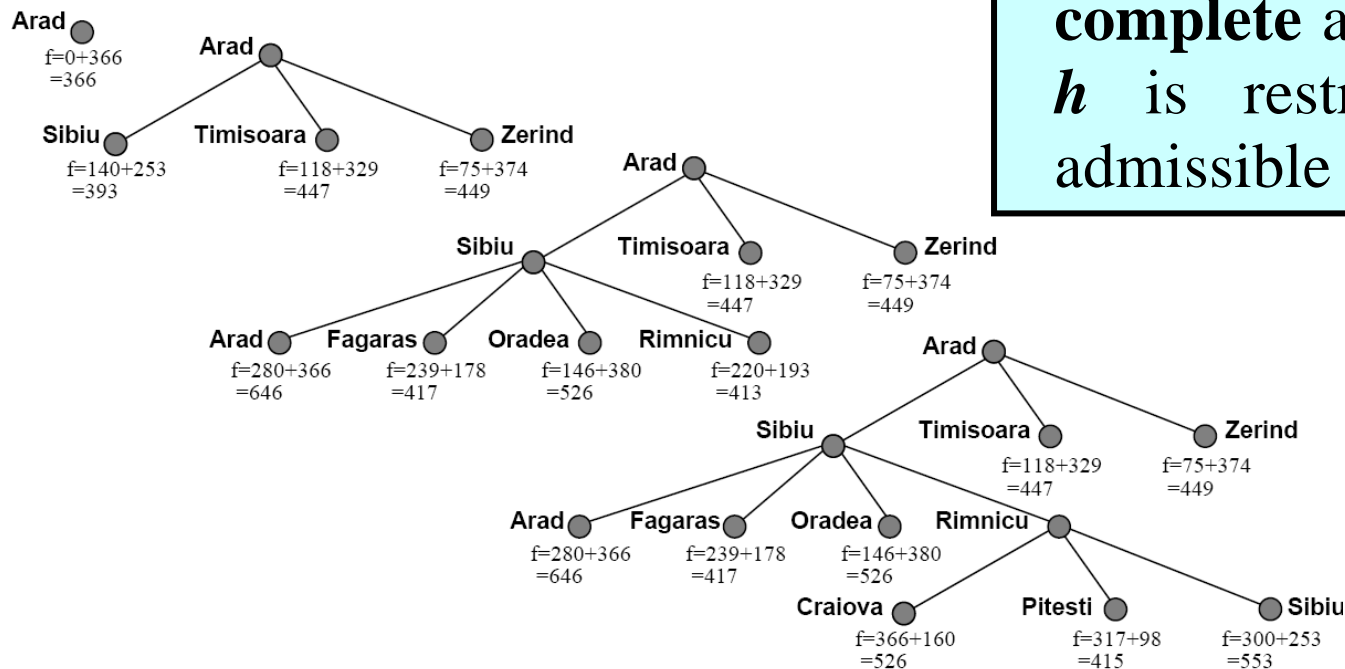
✧ The code is :

**function** A\* SEARCH(*problem*) **returns** a solution or failure.  
**return** BEST-FIRST-SEARCH(*problem*,  $g+h$ )

✧ An example :

✧ **Complete and Optimal:**


A\* search approaches is **complete** and **optimal** if  $h$  is restricted to an admissible heuristic.





## ※ Conceptually prove A\* is optimal and complete

- If  $f$  never decreases along any path out from root, conceptual **contours** can be drawn.



- IF  $f^*$  is cost of optimal solution THEN:

$A^*$  expands all nodes  $f(n) < f^*$

$A^*$  expands some nodes  $f(n) = f^*$

$A^*$  expands no nodes  $f(n) > f^*$

- $A^*$  search is **optimally efficient** for any given heuristic function, i.e. no other optimal algorithm is guaranteed to expand fewer nodes than  $A^*$

✂ **Time complexity**  $O(b^m)$  (bad news)

- Time has exponential growth unless **error** in heuristic function grows no faster than the *log* of the **actual path cost**:

$$| h(n) - h^*(n) | \leq O(\log h^*(n))$$

- For almost all heuristics in practical use, the error is at least proportional to path cost, and therefore results the exponential growth in the time complexity.


## ※Space complexity $O(b^m)$

Because A\* search keeps all generated nodes in memory, empirically, it runs out of space before time.

## ※The behavior of A\* search

- If  $h$  is monotonically increasing, i.e. if it exhibits **monotonicity**,  $f$ -cost will never decrease along any path.
- Almost all **admissible heuristics** are **monotonic**.
- When  $f$  has a non-monotonically increasing  $h(n)$ , it can be fixed by using a **pathmax** equation:

$$f(n') = \max( f(\text{parent}(n')), g(n') + h(n') )$$



$$\text{maybe } f(n) \geq f(n')$$

## Heuristic Functions

Using the 8-puzzle problem as an example to show how a heuristic function can be established.

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

- ✧ Typical solution is **20** steps.
- ✧ Average branching factor is **3**.
- ✧ Blind search would look at  $3^{20} = 3.5 \times 10^9$  states.
- ✧ By checking repeated states would reduce to still **9!=362,880**, Not to mention the memory needed.

✧ Needs an admissible heuristic function to find the shortest solutions. Two candidates:

1.  $h_1$  = the number of tiles that are in the wrong position.
2.  $h_2$  = the sum of the distances of the tiles from their goal position.  
This must be the **city block distance** or **Manhattan distances**, which is the sum of the *horizontal* and *vertical* distances.

✧ For figure 7, the  $h$  values for the start states are:

$$h_1 = 8 \quad h_2 = 2+3+3+2+4+2+0+2$$

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

## ◆ The effect of heuristic accuracy on performance

The quality of a heuristic can be judged by the **effective branching factor  $b^*$**

**$b^*$** : the **branching factor** for an uniform tree of depth  **$d$**  to contain  **$N$**  nodes -- the total number of nodes expanded by  $A^*$ .  
$$N = 1 + b^* + (b^*)^2 + \dots + (b^*)^d.$$

- ✧ A well designed heuristic would have a value of  $b^*$  close to **1**.
- ✧ Compare the result

	Search Cost			Effective Branching Factor		
$d$	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	364404	227	73	2.78	1.42	1.24
14	3473941	539	113	2.83	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Figure 8: Comparison of the search costs and effective branching factors for the ITERATIVE-DEPEENING-SEARCH and  $A^*$  algorithms with  $h_1, h_2$ . Data are averaged over 100 instances of the 8-puzzle, for various solution length.

✧ Will  $h_2$  be always better than  $h_1$ ? Yes

If for any node  $n$ ,  $h_2(n) \geq h_1(n)$ , then  $h_2$  **dominate**  $h_1$ , and  $A^*$  uses  $h_2$  will expand *fewer* nodes.

Every node  $n$  is expanded if

$$f(n) \leq f^* \quad \text{i.e.} \quad h(n) < f^* - g(n)$$

This also means that every node expanded by  $h_2$  will be expanded by  $h_1$  too.

✧ It is always better to use a heuristic function with higher values, as long as it does not overestimate.



## ❖ Inventing heuristic functions

✂ **Relaxed problem:** A problem with fewer restrictions on the operators is called a **relaxed problem**.

*The cost of an exact solution to a relaxed problem is sometimes a good heuristic for the original problem.*

● For example, the 8-puzzle operators:

A tile can move from square A to square B if A is adjacent to B and is blank.

can be reduced to :

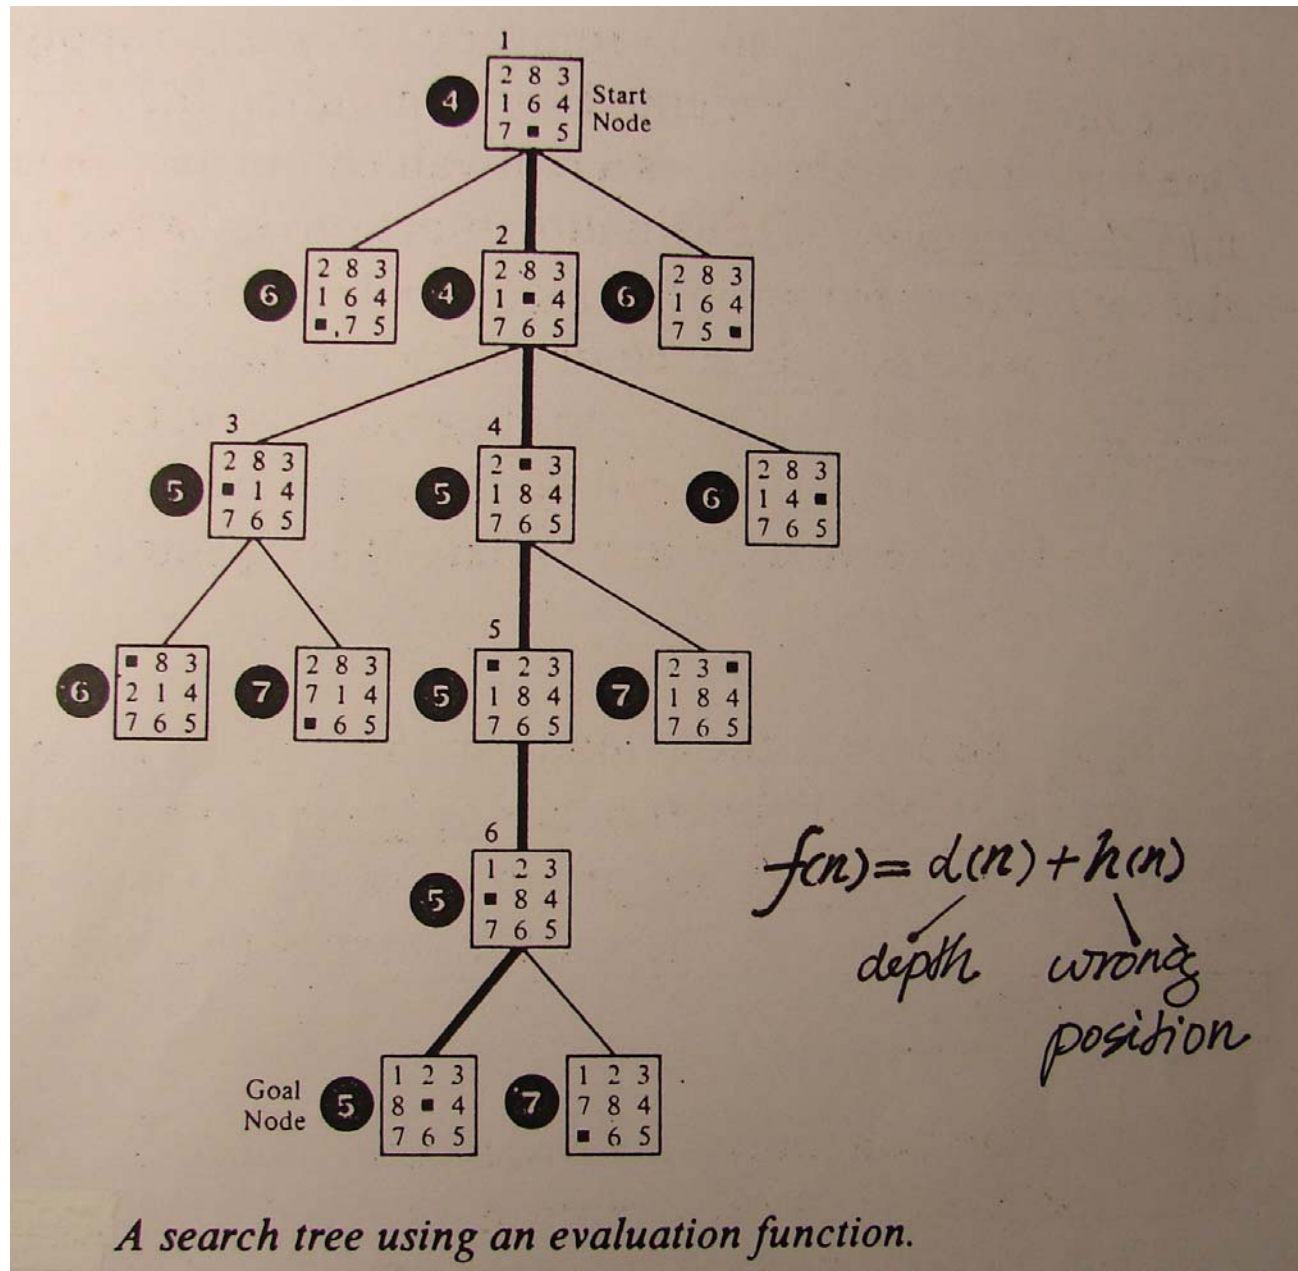
- (a) A tile can move from square A to square B if A is adjacent to B.
- (b) A tile can move from square A to square B if B is blank.
- (c) A tile can move from square A to square B.

## ✂ May use maximum of a set of heuristics:

If a set of admissible heuristics  $h_1, h_2, \dots, h_m$  is available for a problem, and none of them dominates any of others, then just use:

$$h(n) = \max( h_1(n), h_2(n), \dots, h_m(n) ).$$

$h$  is **admissible**, and **dominate** all of the individual heuristics.





## Memory Bounded Search

Available memory space is the first thing to consider when facing complex problem.

### ❖ Iterative deepening A\* search (IDA\*)

Each iteration is a **Depth-first search** that expands all nodes of cost below some *f*-cost limit.

✧ Each iteration expands all nodes inside the **contour** for the current *f*-cost, generating a *new f*-cost for the next iteration.

✧ **Complete:** Yes

✧ **Optimal:** Yes

✧ **Space:**  $O(\text{branching factor} * \text{depth})$

✧ **Time:** similar to  $A^*$ , but depends on the number of different values the heuristic function can take.

1. Typically  $f$  value only increases **two** or **three** times along any solution path -- thus, IDA\* only goes through **two** or **three** iterations
2. Less overhead than  $A^*$  because no priority queue.

**function** IDA\*(*problem*) **returns** a solution sequence

**inputs:** *problem*, a problem

**static:** *f-limit*, the current *f*- COST limit

*root*, a node

*root*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

*f-limit*  $\leftarrow$  *f*- COST(*root*)

**loop do**

*solution, f-limit*  $\leftarrow$  DFS-CONTOUR(*root, f-limit*)

**if** *solution* is non-null **then return** *solution*

**if** *f-limit* =  $\infty$  **then return** failure; **end**

---

**function** DFS-CONTOUR(*node, f-limit*) **returns** a solution sequence and a new *f*- COST limit

**inputs:** *node*, a node

*f-limit*, the current *f*- COST limit

**static:** *next-f*, the *f*- COST limit for the next contour, initially  $\infty$

**if** *f*- COST[*node*] > *f-limit* **then return** null, *f*- COST[*node*]

**if** GOAL-TEST[*problem*](STATE[*node*]) **then return** *node, f-limit*

**for each** node *s* **in** SUCCESSORS(*node*) **do**

*solution, new-f*  $\leftarrow$  DFS-CONTOUR(*s, f-limit*)

**if** *solution* is non-null **then return** *solution, f-limit*

*next-f*  $\leftarrow$  MIN(*next-f, new-f*); **end**

**return** null, *next-f*

## ✧ Problems in more complex domains:

1. When each contour expands very few (maybe just 1) state.

Example domain: In the Traveling Salesman problem, the heuristic value is different for every state. Thus, each iteration will expand exactly **one more** state.

If  $A^*$  expands  $N$  nodes, IDA\* will expand  
 $1+2+\dots+N = O(N^2)$  nodes.

**A possible solution:  $\epsilon$ -admissible algorithm** Increase the  $f$ -cost limit by at least some  $\epsilon$  on each iteration.

This, however, finds solutions that can be worse than the true optimal solution by at most  $\epsilon$ .

2. No memory between searches: may repeat searching from an already seen node (solution: see SMA\*).

## ◆ Simplified Memory Bounded A\* (SMA\*)

Try to use **any available memory** to carry out the search, especially to avoid repeated states.

- ✧ SMA\* avoids repeated states if record of them fits in memory.
- ✧ In generating a successor, if there is no memory left, the node with the high  $f$ -cost in the queue is dropped.
- ✧ The quality of the best path in a forgotten sub-tree is remembered in a node. In case of regenerating, only the path is generated.
- ✧ Simplified SMA\* algorithm is in *Figure 12*.



- ✧ **Complete:** Yes, if the available memory can hold the shallowest solution path.
- ✧ **Optimal:** Yes, if the available memory can hold the shallowest optimal solution path.
- ✧ Search is **optimally efficient** for the entire search tree that fits in the available memory.

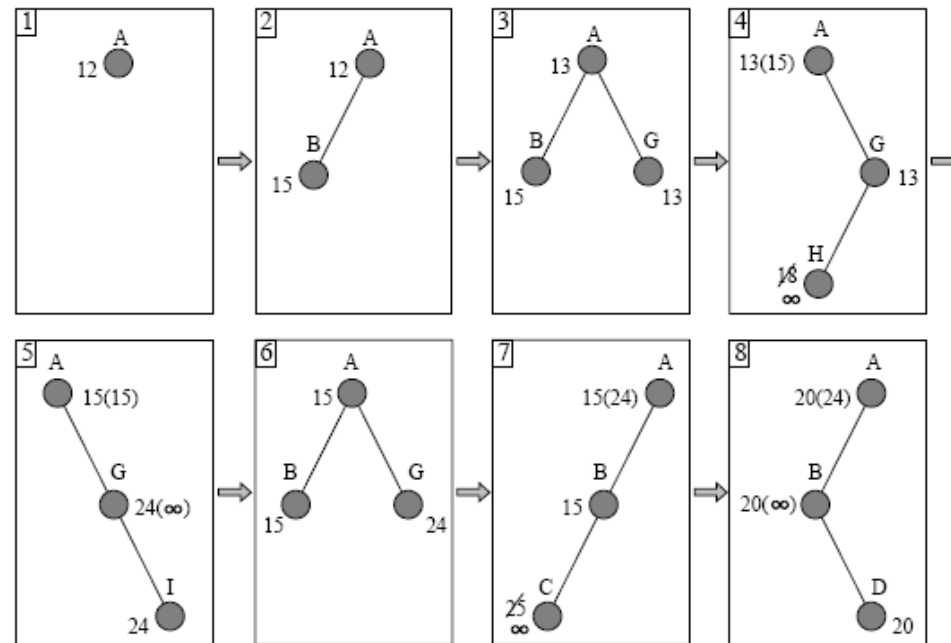



Figure 11: Progress of an SMA\* search with a memory size of three nodes, on the state space shown at the top. Each node is labelled with its *current*  $f$ -cost. Values in parentheses show the value of the best forgotten descendant.

```

function SMA*(problem) returns a solution sequence
  inputs: problem, a problem
  static: Queue, a queue of nodes ordered by f-cost

  Queue  $\leftarrow$  MAKE-QUEUE({MAKE-NODE(INITIAL-STATE[problem])})
  loop do
    if Queue is empty then return failure
    n  $\leftarrow$  deepest least-f-cost node in Queue
    if GOAL-TEST(n) then return success
    s  $\leftarrow$  NEXT-SUCCESSOR(n)
    if s is not a goal and is at maximum depth then
      f(s)  $\leftarrow \infty$ 
    else
      f(s)  $\leftarrow$  MAX(f(n), g(s)+h(s))
    if all of n's successors have been generated then
      update n's f-cost and those of its ancestors if necessary
    if SUCCESSORS(n) all in memory then remove n from Queue
    if memory is full then
      delete shallowest, highest-f-cost node in Queue
      remove it from its parent's successor list
      insert its parent on Queue if necessary
    insert s on Queue
  end

```

Figure 12: Sketch of the SMA\* algorithm. Note that numerous details have been omitted in the interests of clarity.

## Iterative Improvement Algorithms

For problems in which the *state description itself contains all the information needed for a solution*, the iterative improvement algorithms provide the most practical approach.

- ✧ The path by which the solution is researched is irrelevant.
- ✧ The general idea is to start with a complete configuration and to make modifications to improve its quality.
- ✧ Consider to lay out all the states on the surface of a landscape as in (*Figure 13*)
- ✧ The idea is to move around the landscape trying to find the highest peaks, which are the optimal solution.
- ✧ Iterative improvement algorithms usually keep track of only the current state.




Figure 13: Iterative improvement algorithms try to find peaks on a surface of states where height is defined by the evaluation function.

## ◆ Hill-climbing ( or gradient descent) search

Always moves in the direction of increasing value. The algorithm does not maintain a search tree.

```
function HILL-CLIMBING(problem) returns a solution state
  inputs: problem, a problem
  static: current, a node
           next, a node

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  loop do
    next  $\leftarrow$  a highest-valued successor of current
    if VALUE[next] < VALUE[current] then return current
    current  $\leftarrow$  next
  end
```

## ✧Problems:

1. **local maxima**: if the algorithm is trapped on a local maxima, it will halt even though the solution is far from satisfactory.
2. **Plateaux**: If the algorithm hits an area that is basically flat, it will do random walk.
3. **Ridge**: (suppose the peak is on the same ridge) there may not be an operator that allows one to change to an adjacent higher state.

✧**Random-restart hill-climbing**, which starts a series of hill-climbing searches from randomly generated initial state, may solve the problem and find the optimal solution given enough iterations.

## Simulated annealing

Jump randomly and allow to take some **downhill** steps (to escape the local maximum) with decreasing frequency (probability).

**function** SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

**inputs:** *problem*, a problem

*schedule*, a mapping from time to “temperature”

**static:** *current*, a node

*next*, a node

*T*, a “temperature” controlling the probability of downward steps

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**for** *t*  $\leftarrow$  1 **to**  $\infty$  **do**

*T*  $\leftarrow$  *schedule*[*t*]

**if** *T*=0 **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  VALUE[*next*] – VALUE[*current*]

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E/T}$



- ✧ The probability to move to a worse state decreases exponentially with the “**badness**” of the move--the amount  $\Delta E$  by which the evaluation is worsened.
- ✧ The probability is also controlled by the **T**, the “**temperature**”, which is determined by a *schedule* as a function of how many cycles already have been completed.
- ✧ Higher T, higher probability to make “**bad**” moves.
- ✧ If the schedule lowers T slowly enough, the algorithm will find a global optimum.

