

Game Playing

Game playing is an idealization of worlds in which hostile agents act so as to diminish one's well-being.



Introduction: Games as Search Problems



Why Game Playing

- Simple Rules
- Fully accessible world means *precise representation* of a game as a search through a space of possible game positions



Characteristics

- The presence of an opponent makes the decision problem somewhat more complicated than the simple search problems.
- The opponent introduces uncertainty- all game playing programs must deal with the contingency problem.
- **Much too hard to solve.** It relies heavily on one's past experience - thus games are much more like the real world than the standard search problems.
- **Time** is very important - must make the best use of time to reach good decisions, when reaching optimal decision is impossible.



Perfect Decisions in Two-Person Games :

A game can be formally defined as a kind of search problem with the following components:

The **initial state** : board position, whose move, etc..

A set of **operators** : define legal moves.

A terminal test: determines when the game is over, i.e. determines the **terminal states**.

A **utility function** (also called the **payoff function**) : gives a numeric value for the outcome of a game.



Minimax Algorithm

✧ The algorithm must find for **MAX** a strategy that will lead to a winning terminal state regardless of what **MIN** does, including the correct move for **MAX** for each possible move by **MIN**.

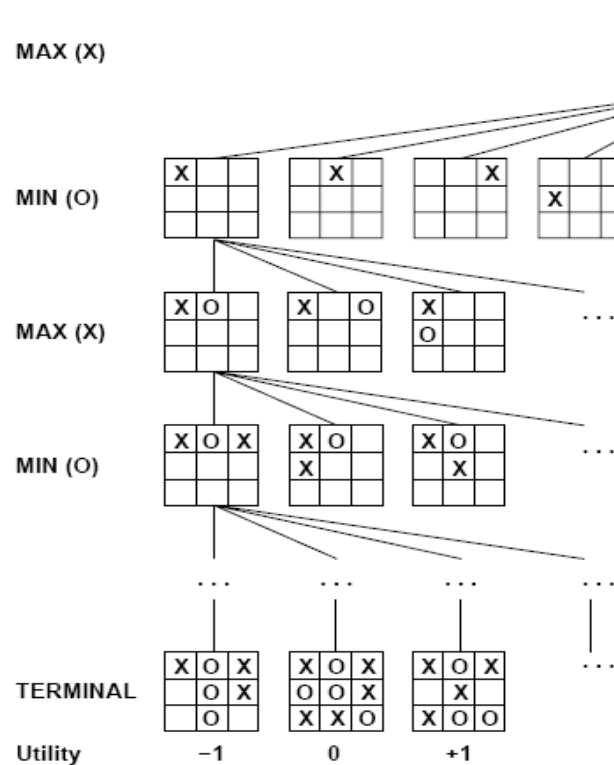


Figure 1: A (partial) search tree for the game of Tic-Tac-Toe.




Figure 2: A two-ply game tree as generated by the minimax algorithm.

✧ The algorithm must do the minimax decision, in which it maximizes the utility value under the assumption that the opponent will play perfectly to minimize it.

✧ The algorithm:

1. Generate the whole game tree.
2. Apply utility function to leaf states to get their utility values
3. Recursively do from the leaves: **min** chooses the min among choices, **max** chooses the max, backing up toward the root, one player at a time.
4. At the top **max** goes to the node that maximizes outcome.

✧ The code:

```

function MINIMAX-DECISION(game) returns an operator

  for each op in OPERATORS[game] do
    VALUE[op] ← MINIMAX-VALUE(APPLY(op, game), game)
  end
  return the op with the highest VALUE[op]



---


function MINIMAX-VALUE(state, game) returns a utility value

  if TERMINAL-TEST[game](state) then
    return UTILITY[game](state)
  else if MAX is to move in state then
    return the highest MINIMAX-VALUE of SUCCESSORS(state)
  else
    return the lowest MINIMAX-VALUE of SUCCESSORS(state)

```

Figure 3: An algorithm for calculating minimax decision.

✧ Finds **optimal** strategy

✧ **Time:** $O(\text{branch_factor}^{\text{depth}})$

✧ **Space:** $O(\text{branch_factor} * \text{depth})$ (if implement as DFS)

✧ **The Problem:** Must assume that the whole tree can be

generated.



Imperfect Decisions:

The algorithm must cut off the search earlier and apply a heuristic **evaluation function** to the leaves of the tree, thus -

Utility function \Rightarrow EVAL function

Terminal test \Rightarrow CUTOFF-TEST



Evaluation Function

Returns an *estimate* of the expected utility of the game from a given position.

✧ Usually use **material values**. For example:




(a) White to move
Fairly even



(b) Black to move
White slightly better



(c) White to move
Black winning



(d) Black to move
White about to lose

Pawn = 1,
knight or bishop = 3,
rook = 5, and
queen = 9, etc

“good pawn structure” and “king safety” = 0.5

Figure 4: Some chess positions and their evaluations.

- ✧ The performance of a game-playing program is extremely dependent on the quality of its evaluation function.
- ✧ Characteristics of a good evaluation function:
 1. Should agree with the real utility function on **terminal states**.
 2. Should be **quick** to compute. is usually a trade-off between the accuracy of the evaluation function and its time cost.
 3. Should accurately reflect the actual chances of winning.

✧Types of evaluation functions:

1. **Weighted linear function:** sum of weights time features

$$w_1f_1+w_2f_2+....+w_nf_n$$

Assume that the value of a piece can be judged independently of the other pieces present on the game board.

2. **Non-linear:** (e.g. neural nets for backgammon)



Cutting Off Search


✧ When to cutoff searching: Using

1. Fixed depth limit - the cutoff test succeeds for all nodes at or below depth *d*.
2. Iterative deepening - when time runs out, the program returns the move selected by the deepest completed search.

- ✧ Problems: May cutoff before some really good move for opponent. For example, in (*Figure 4.d*), it seems white is winning but actually its queen is losing.
- ✧ The Solution: Continue search down the move until **quiescence** positions are reached.

A quiescent position is a position that is impossible to exhibit wild swing, i.e. neither player gains much. This kind of search is called the **Quiescence search**.

✧ Another problem: The **horizon problem**: when opponent is about to gain, program may try to stall the inevitable by pushing that gain beyond its horizon.



Black to move

Figure 5: The horizon problem. A series of checks by the black rook forces the inevitable queening move by white “over the horizon” and makes this position look like a slight advantage for black, when it is really a sure win for white.

The problem with fixed-depth search is that it believes that these stalling moves have avoided the damaging move by the opponent.



Alpha-Beta Pruning:

In reality, under time constraint, it is impossible to search many ply if every node in the search tree is examined.

For example: Can search 1000 position a second.

With branching factor 35,

150 seconds per move can do only 3-4 ply only.



The alpha-beta pruning algorithm

Can still compute the correct minimax decision after eliminating a branch of the search tree.

Using

alpha α = best choice so far at any choice point along the path for MAX, and
beta β = best choice(i.e. the lowest-value) so far for Min.

to prun a subtree (i.e. terminate the recursive call) as soon as it is known to be worse than the current α or β value





Figure 6: The two-ply games tree as generated by alpha-beta.



function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\text{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Figure 7: Alpha-beta pruning: the general case. If *m* is better than *n* for Player, we will never get to *n* in play.



Effectiveness of alpha-beta pruning

- ✧ Effectiveness depends on the order successors are examined.
- ✧ If the most likely to be the best successor nodes can be examined first(?) then the algorithm needs $O(\text{branch}^{(\text{depth}/2)})$ to pick the best move instead of $O(\text{branch}^{\text{depth}})$ of minimax.
- ✧ Effective branch factor is $\sqrt{\text{branch}}$ and the alpha-beta would be able to look **twice** as far in same time.

Randomly pick next successor:

Asymptotic complexity $O((b/\log b)^{\text{depth}})$ (approx $O(b^{\text{depth}})$). Not good enough!

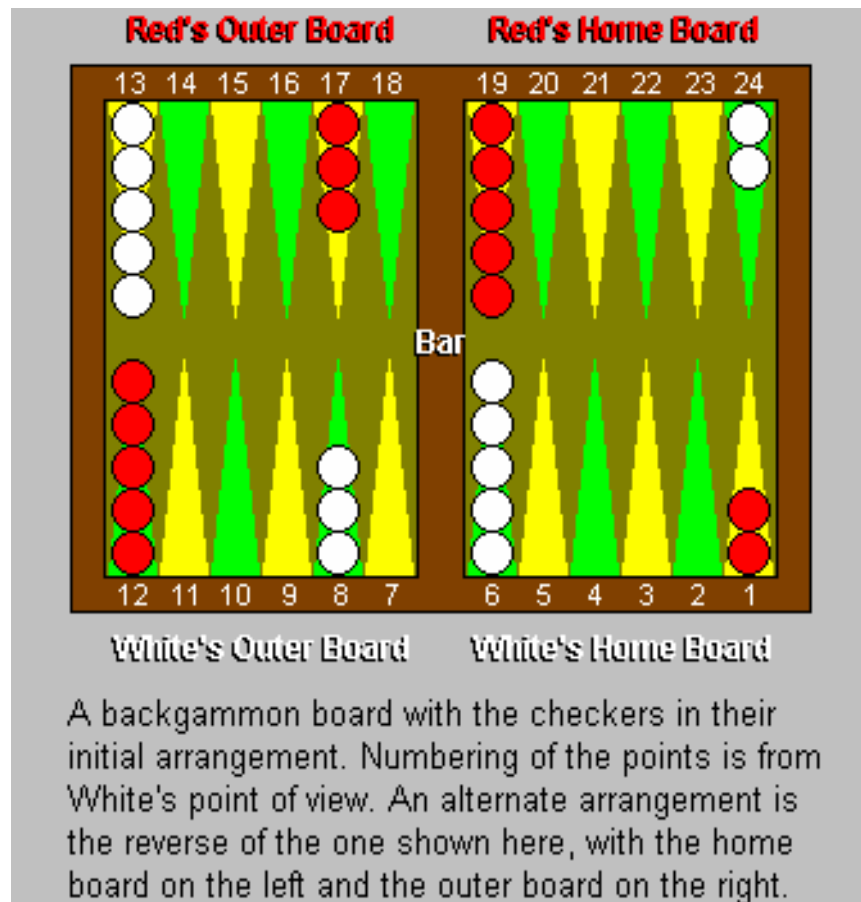
✧Improvements that can be achieved easily:

- 1.Use even a simple ordering scheme (e.g. try to capture first, then threats, then forward moves,...)
- 2.Do iterative deepening and determine the ordering of successors by values of previous iteration.



Games that include an element of chance:

Games that incorporate luck and skill introduce the **unpredictability** by including a random element such as throwing dice.



Chance nodes: nodes to denote the die roll or whatever associated to random elements.

A game with a random element must include in its game tree chance nodes in addition to **MAX** and **MIN** nodes.




Figure 9: A typical backgammon position. The aim of the game is to move all one's pieces off the board. White moves clockwise toward 25, and black moves counterclockwise toward 0. A pieces can move to any position except one where there are two or more of the opponent's pieces. If it moves to a position with one opponent piece, that piece is captured and has to start its journey again from the beginning. In this position, white has just rolled 6-5 and has the four legal moves: (5-10, 5-11), (5-11, 19-24), (5-10, 10-16), and (5-11, 11-16).

Each possible position no longer has a **definite minimax value**, but an **average** or **expected value** which is taken over all the possible dice rolls that could occur.



Expectimax and Expectimin values

- ✧ For terminal nodes, the expected value is calculated by using utility function, just like in deterministic games.
- ✧ At a **chance node**(for MAX), compute the expected value **expectimax** of each move by summing the probability of a state by its value.

$$\mathbf{expectimax}(C) = \sum_i P(d_i) \max_{s \in S(C, d_i)} (Utility(s))$$

where

$P(d_i)$ is the chance or probability of obtaining a roll.

$S(C, d_i)$ is the set of positions generated by applying the legal moves for dice roll $P(d_i)$ to the position at C .

```

expectimax = 0;
for each dice roll  $d_i$  from chance node C:
{
  s = state with max utility under  $C, d_i$ 
  expectimax +=  $P(d_i) * \text{utility}(s)$ 
}

```

- ✧ At a **chance node**(for MIN), compute the expected value **expectimin** of each move by summing the probability of a state by its value using a formula that is analogous to expectimax.

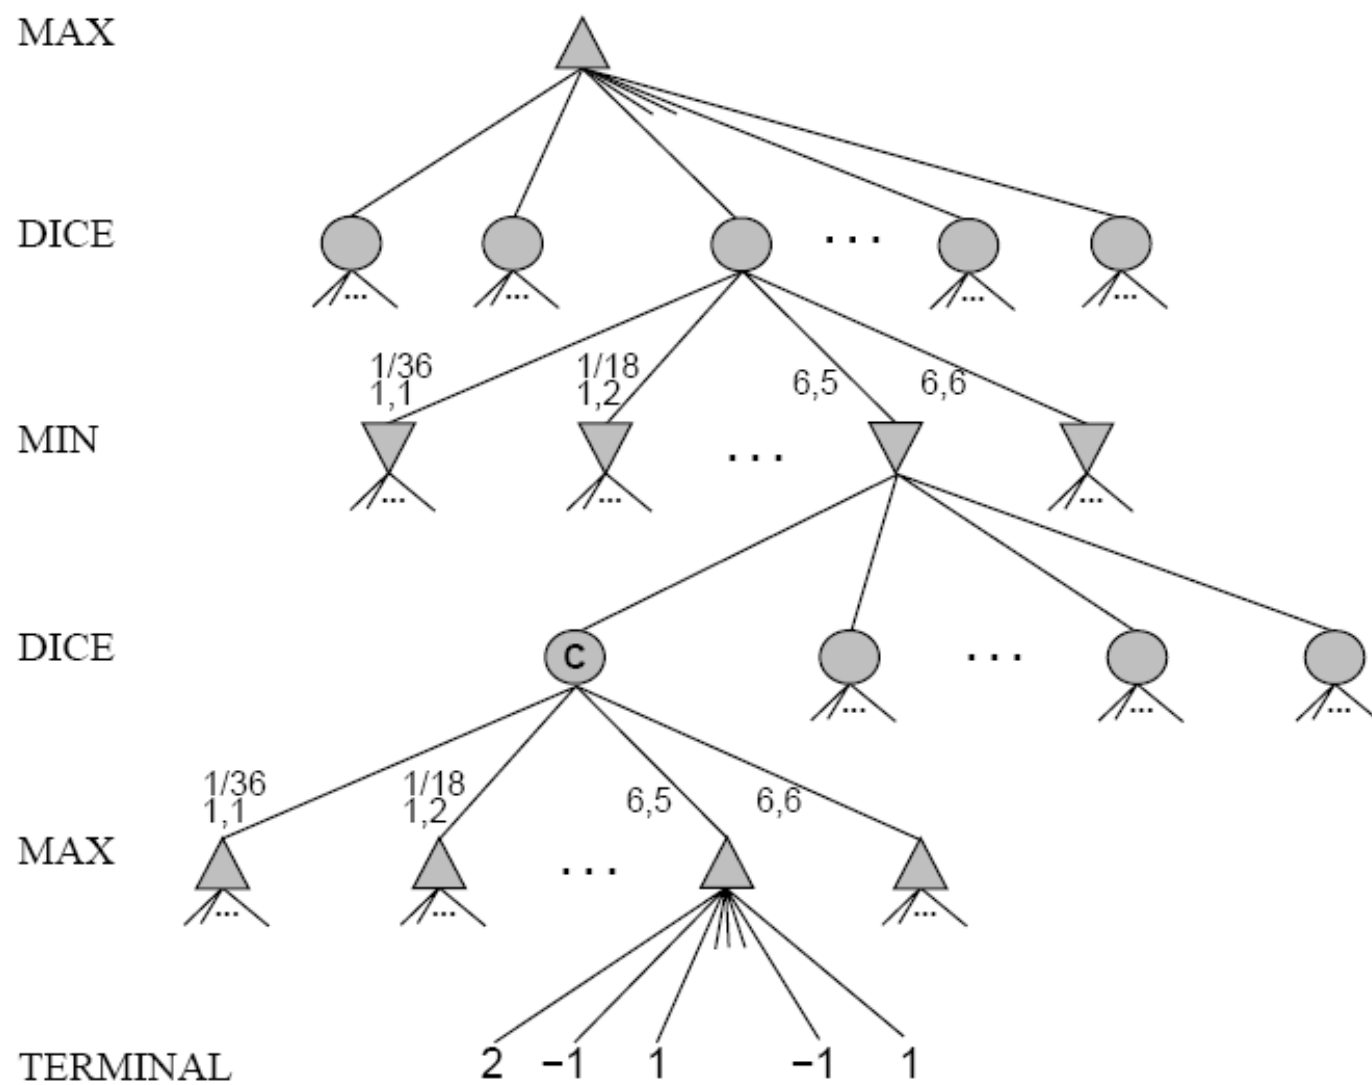


Figure 10: Schematic game tree for a backgammon position.

❖ Position evaluation in games with chance nodes

- ✧ The evaluation function must be a *positive linear* transformation of the likelihood of winning from a position.
- ✧ This is to avoid the situations in *Figure 5.11*




Figure 11: An order-preserving transformation on leaf values changes the best move.



Complexity of Expectiminmax

✧ **Complexity: $O(\text{branch}^{\text{depth}} * \text{num_rolls}^{\text{depth}})$**

✧ IF there are limits on highest and lowest payoffs THEN we can prune.



State-of-the-Art Game Programs:

✧ **Chess:** Deep Thought, Deep Blue. (*Figure 12*)

✧ **Checkers:**

1. Samuel's Checker Player: an original game player that learned with a linear evaluation function.

2. Chinook: Better than all but Dr. Tinsley.

✧ **Othello:** Computers are better than humans.

✧ **Backgammon:** Gerry Tesauro's program is reliably ranked among the top three players in the world.

✧ **Go:** Not yet, branch factor is too high.






Figure 12: Ratings of human and machine chess champions.



■ Deep Blue

- Beaten by Garry Kasparov in 1996. World Champion Garry Kasparov took on Deep Blue computer in a 6 game match. He lost game 1, won game 2, drew games 3 and 4, and won games 5 and 6 to win the match with a 4-2 score.
- Defeated Garry Kasparov in 1997 match (2 wins 3 draws 1 lose)
- The sixth and final game in the series ended in a draw after 27 moves, leaving both sides with three points each, having each won a game with three other matches drawn. (Second match, Sun. February 09, 2003)
- The third match was in New York, on 26 January 2003 Garry Kasparov vs. Deep Junior. Draw.
- 32 P2SC Processors, capable of searching 50 to 100 billion positions within three minutes
- 1000 times faster than its predecessor, Deep Thought

■ Pacific Blue

- Powered by 8,192 IBM Power3 processors and has 160 terabytes of disc storage.
- Process information at 12.3 trillion calculations per second.
- Being used for nuclear weapon testing.

■ **Blue Gene**

- Running on 32 specially constructed CPU chips containing both memory and communication circuits.
- Possible to do more than a quadrillion calculations per second.
- Blue Gene will be 500 times faster than Pacific Blue and a thousand times faster than Deep Blue
- Blue Gene will study gene sequences and the complex shapes of human proteins
- Also be used to study weather and global climate changes,



Discussion:

❖ Minimax + Alpha-beta : The only solution?

- ✧ Minimax is an optimal method for selecting a move from a given search tree *provided the leaf node evaluations are exactly correct.*
- ✧ In reality, it might not be the case, see (*Figure 13*).
- ✧ It might have to combine the evaluation with the *probability distribution.*

❖ Meta-reasoning (reasoning about reasoning)

- ✧ The idea of the *utility of a node expansion* - good search algorithm should select node expansions of high utility.
- ✧ If there are no node expansions whose utility is higher than their cost(in terms of time), then the algorithm should stop searching and make a move.
- ✧ Alpha-beta algorithm is designed to calculate the values of *all the legal moves*- even if there is only **one** legal move.

Goal-directed reasoning or planning

✧ The actual reasoning method of human player.